


Accuracy vs. inference latency

GPU (T4)



CPU (single-thread)

